



Cairo University  
Faculty of Engineering  
Computer Engineering Department

## CMPS458 Reinforcement Learning Report Assignment 1 Policy Iteration

Team Name/Number:  
Mohamed Ahmed Ibrahim Sobh 1210288

Omar Ahmed Ibrahim 1210020

*Supervisor:* Ayman AboElhassan

October 23, 2025

# Deliverables

Repo link: [Github Repository](#)

Video record link: add link here

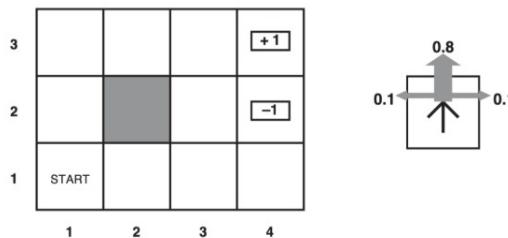


Figure 1: Example grid-world environment used for testing.

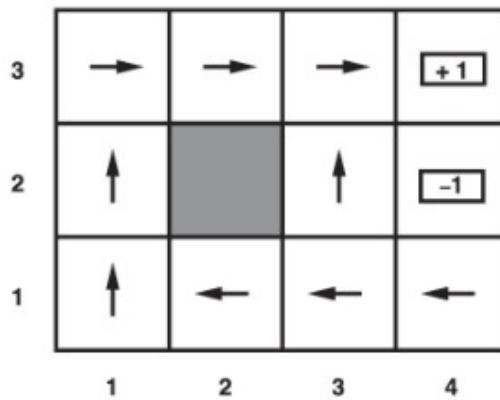


Figure 2: Example of a learned policy after convergence.

Table 1: Core parameters used in the policy iteration experiments.

Parameter	Symbol	Default Value	Description
Discount factor	$\gamma$	0.9	Future reward weighting
Convergence threshold	$\theta$	$1 \times 10^{-6}$	Stopping criterion for evaluation
Maximum iterations	$N_{iter}^{max}$	1000	Upper limit on policy updates
Step reward	$R_{step}$	-1	Penalty for each move
Goal reward	$R_{goal}$	10	Reward for reaching the goal state
Bad cell penalty	$R_{bad}$	-10	Penalty for entering a bad cell

# Discussion

## 0.1 Experiments

Table 2: Effect of varying the discount factor  $\gamma$  on policy iteration results (seed=40501, maze=5, bad=2, correct hit=0.7 BUT the simulation steps are based on policy only).

$\gamma$	Observation	Outcome / Behavior
0.3	Agent focuses on immediate rewards	5 Policy Iterations AND 7 Simulation Steps
0.9	Balances short and long-term rewards	4 Policy Iterations AND 5 Simulation Steps

Table 3: Effect of convergence threshold  $\theta$  on stability and convergence (maze=100, bad=40).

$\theta$	Observation	Outcome / Behavior
$10^{-1}$	Fast convergence but less accurate	116 Iters to reach Stable Policy, Time taken = 15.75s
$10^{-6}$	Slower convergence with high precision	83 Iters to reach Stable Policy, Time taken = 19.57s

Table 4: Effect of changing reward values on agent behavior and learning outcome (maze=40, bads=20).

Reward Configuration	Observation	Expected Outcome
Goal=5, Bad=-1, Step=-0.1	Small penalty difference causes uncertain exploration	25 Policy Steps AND 55 Simulation Steps
Goal=10, Bad=-10, Step=-1	Balanced reward setup encourages stable learning	15 Policy Steps AND 55 Simulation Steps
Goal=20, Bad=-50, Step=-1	Strong contrast between rewards and penalties	20 Policy Steps AND 57 Simulation Steps

## 0.2 Question Answers

### 1. What is the state-space size of the 5x5 Grid Maze problem?

The state space size depends on how we define the state. In our Grid Maze environment, the state consists of Agent, Goal, Bad cell 1, and Bad cell 2 coordinates. For a  $5 \times 5$  grid, we have 25 possible positions. Since we have 4 unique positions, we can say that the state space size is  $25 \times 24 \times 23 \times 22 = 303,600$ . If we fix the goal and bad cell positions then the state space size is 25 states only

### 2. How to optimize the policy iteration for the Grid Maze problem?

#### (a) Modified Policy Iteration

- Do not run policy evaluation to full convergence - Use a fixed number of evaluation sweeps (e.g.,  $k = 3$ ) before improvement - This balances evaluation accuracy with computational speed

#### (b) Early Termination

- Set a practical threshold ( $\theta$ ) for policy evaluation convergence - Stop evaluation when changes are negligible (e.g.,  $\theta=1e-4$ )

#### (c) Exploiting Problem Structure

- Use Manhattan distance heuristics to initialize value function - Initialize policy to point toward the goal - Skip evaluation for terminal states (goal and bad cells)

#### (d) Asynchronous Updates

- Update states in-place during policy evaluation instead of using a separate copy
- Prioritize updates for states more likely to be visited - Focus computation on the "critical path" from start to goal

#### (e) Parallel Processing

- Evaluate multiple states simultaneously using vectorization - NumPy operations can process the entire grid at once

### 3. How many iterations did it take to converge on a stable policy for 5x5 maze?

The exact number depends on:

- (a) **Initial policy:** Random initialization may take longer
- (b) **Maze configuration:** Distance from start to goal affects convergence
- (c) **Discount factor ( $\gamma$ ):** Higher values (closer to 1.0) may require more iterations
- (d) **Convergence threshold ( $\theta$ ):** Stricter thresholds require more iterations

Using seed 40501, 10, -1, -10 reward system, 0.9 discount factor, and convergence factor 1e-6. Convergence was achieved in **4 iterations** of policy evaluation and improvement cycles.

### 4. Explain, with an example, how policy iteration behaves with multiple goal cells.

Policy iteration can handle multiple goal cells by forming a **multi-objective optimization problem**, where the agent learns to reach the **nearest** or **most valuable** goal.

**Scenario:**  $5 \times 5$  grid with two goal cells ( $G_1$  and  $G_2$ ):

```

G1 . . . .
. X . X .
. . S . .
. . . .
. . . G2

```

### 1. Modified Reward Function:

```

1 def _get_reward(self, pos):
2     if pos in self.goal_positions: # instead of pos == self.goal_pos
3         return 100.0
4     elif pos in self.bad_cells:
5         return -100.0
6     else:
7         return -1.0

```

### 2. Value Function Behavior:

- The value function develops **multiple peaks** (one at each goal).
- States closer to any goal have higher values.
- The optimal policy routes agents to the **nearest goal**.

### 3. Policy Convergence:

- The grid divides into **regions of attraction**.
- Each region's policy points toward the nearest goal.
- The boundary between regions depends on distance, obstacles, and transition probabilities.

**Example (Equal Rewards):** Both  $G_1$  and  $G_2$  give +100 reward.

```

G1 → → ↓ ↓
↑ → → ↓ ↓
↑ ↑ S ↓ ↓
↑ ↑ ← ↓ ↓
↑ ← ← ← G2

```

*Observations:*

- Top-left states move toward  $G_1$ , bottom-right toward  $G_2$ .
- A diagonal boundary separates the two regions.
- The agent always takes the shortest path.

**Example (Different Rewards):** If  $G_1 = +100$  and  $G_2 = +200$ :

- More states route toward  $G_2$ .
- The boundary shifts toward  $G_1$ .
- Only nearby states prefer  $G_1$ .

### **Mathematical Explanation:**

For each state  $s$ , policy improvement follows:

$$\pi(s) = \arg \max_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V(s')]$$

With multiple goals,  $V(s')$  reflects the best reachable goal's return, so the policy naturally guides the agent to the **optimal goal**.

**5. Can policy iteration work on a 10x10 maze? Explain why.**

Yes, policy iteration can work on a  $10 \times 10$  maze, but the computational cost grows quadratically with the number of states (100 states  $\rightarrow$  400 state-action pairs). It remains feasible, but slower.

**6. Can policy iteration work on a continuous-space maze? Explain why.**

No, standard policy iteration assumes a discrete and finite state space. In a continuous environment, the number of possible states is infinite, making it impossible to represent the value function as a lookup table.

**7. Can policy iteration work with moving bad cells (like Pac-Man moving ghosts)? Explain why.**

Not directly. Policy iteration assumes a stationary environment (fixed transition probabilities). If bad cells move, the environment becomes *non-stationary*.